

Native or Web? A Preliminary Study on the Energy Consumption of Android Development Models

Wellington Oliveira, Wesley Torres, Fernando Castor, Bianca H. Ximenes
Federal University of Pernambuco
Recife, PE, Brazil
{woj, wst, castor, bhxmm}@cin.ufpe.br

Abstract—Energy consumption has become an increasingly important topic in software development, especially due to the ubiquity of mobile devices, and the choice of programming language can directly impact battery life. This paper presents a study on the issue of energy efficiency on the Android platform, comparing the performance and energy consumption of 33 different benchmarks in the two main programming languages employed in Android development: Java and JavaScript. Preliminary results show that Java applications may consume up to 36.27x more energy, with a median of 2.28x, than their JavaScript counterparts, in benchmarks that are mostly CPU-intensive. In some scenarios, though, Java benchmarks exhibited better energy efficiency, with JavaScript consuming up to 2.27x more energy. Based on these results, two Java applications were re-engineered, and through the insertion of JavaScript functions, hybrid applications were produced. In both cases, improvements in energy efficiency were obtained. Considering that Android apps written in Java are the norm, results from this study indicate that using a combination of JavaScript and Java may lead to a non-negligible improvement in energy efficiency.

I. INTRODUCTION

As smartphones become popular and diverse platforms emerge (Android, iOS, Windows Phone), developers must choose whether to create applications using the native language of a mobile OS or the Web toolkit (HTML, CSS and JavaScript), subsequently porting the app using a specific framework (Cordova¹, Ionic²). In addition to having to decide which model to employ, developers have little to no information available on the difference between the performance and battery consumption for these approaches, making it more difficult to determine which one is more adequate.

Currently, the most popular OS in smartphones and tablets is Android³. Previous work has analyzed the energy consumption of Android from a number of different perspectives [13], [6], [8]. However, to the best of our knowledge, the impact of different approaches for Android development on energy consumption has not yet been investigated. Android apps can be written entirely in Java (**native** apps), in JavaScript-related technologies (**Web** apps), or in a combination of both (**hybrid** apps). Most of the Android apps, however, are written entirely in Java. In a sample of 109 projects we examined from F-Droid⁴, only 4% use JavaScript. It is not yet clear whether this approach leads to energy-efficient applications.

This paper sheds some light on the issue of the energy efficiency of Android app development approaches. We compare energy consumption and performance of 33 benchmarks developed by several authors from Rosetta Code⁵ and The Computer Language Benchmark Game⁶. To measure energy consumption, we used Android's Project Volta [1]. Our goal with this study is to provide a preliminary answer to the following research question:

- **RQ1. Is there a more energy-efficient approach among the two most common Android development models?**

JavaScript versions of 26 out of the 33 analyzed benchmarks exhibited lower energy consumption. The Java versions of six of these benchmarks outperformed their JavaScript counterparts, even though they consumed more energy. This result indicates that, at least for CPU-intensive apps, Java may not be the most energy-efficient solution. Most of their execution time is spent waiting for user input or using sensors. This led us to question whether one could save energy by using a hybrid approach and, e.g., adopting JavaScript in the more CPU-intensive parts of applications. Thus, we also provide an initial answer for the following additional research question:

- **RQ2. Is it possible to reduce the energy consumption of an app built using a single approach by making it hybrid?**

We reengineered two existing apps from F-Droid written in Java and made parts of them run in JavaScript. We analyzed different models for the Java part of the apps to invoke the JavaScript part and measured the energy consumption in all the cases. Our results indicate that it is possible to save energy using this hybrid approach. In one of the apps, TriRose, the hybrid version saved up to 30% energy by grouping invocations to the JavaScript part. Performing the same grouping of operations using only Java yielded only marginal gains in terms of energy consumption.

Knowing whether small modifications in the code promote a non-negligible reduction in energy consumption empowers developers. Moreover, tool builders can introduce cross-language refactorings that support developers in reengineering existing applications when a hybrid approach may be beneficial.

¹<https://cordova.apache.org/>

²<http://ionicframework.com/>

³<http://www.idc.com/prod/serv/smartphone-os-market-share.jsp>

⁴<http://f-droid.org>

⁵http://rosettacode.org/wiki/Rosetta_Code

⁶<http://benchmarksgame.alioth.debian.org>

TABLE I
THE SELECTED SET OF BENCHMARKS AND APPLICATIONS. ALL OF THEM
INCLUDE VERSIONS IN BOTH JAVA AND JAVASCRIPT.

Source	Benchmark or App
Rosetta Code	SeqNonSquares, Perfect Number, HofstadterQ, Zero-One Knapsack, Knapsack Unbounded, nQueens, MergeSort, Matrix Multiplication, GnomeSort, Sieve of Eratosthenes, Man or Boy, BubbleSort, CountingSort, QuickSort, Knapsack Bounded, Happy Numbers, HeapSort, InsertSort, Tower of Hanoi, Count in factors, Combinations, PancakeSort and ShellSort
The Computer Language Benchmark Game	BinaryTree, Fannkuch, Fasta S.Core, Fasta M.Core, Nbody, Spectral, RegenDna S.Core, RegenDna M.Core, RevComp and Knuceotide
F-Droid	anDOF and Tri Rose

II. METHODOLOGY

The aim of this study is to analyze the two most popular programming approaches for Android app development and to establish whether they differ in terms of energy efficiency and performance. In this Section, we describe how we selected the analyzed benchmarks (Section II-A) and explain our experimental procedure (Section II-B).

A. Benchmarks

Benchmarks were extracted from Rosetta Code and The Computer Language Benchmark Game (TCLBG). Rosetta Code is a programming chrestomathy site. It includes a large number of programming tasks and solutions to these tasks in different programming languages. TCLBG is a website whose main purpose is to compare the performance of several programming languages. Both have been employed in prior work for comparing different programming languages [9] and to analyze energy efficiency [7].

The benchmark set encompasses 23 benchmarks from Rosetta Code and 10 from TCLBG. Originally, all benchmarks had versions written in both Java and JavaScript. Table I lists all the benchmarks analyzed in this study. Most of the benchmarks from Rosetta have already been used in other studies [9]. All benchmarks from TCLBG were used.

Since the benchmarks from Rosetta Code were not built with optimal performance in mind, their performances vary widely. For example, in the benchmark nQueens, the solutions available at Rosetta Code took 20s to finish in Java and 69s in JavaScript. By converting the JavaScript version to use the same algorithm as the Java version, it took 12s to finish. As the research focus is on the languages and not algorithm implementations, we analyzed all of Rosetta Code's benchmarks and, whenever necessary, performed similar modifications. Those modifications make comparisons more balanced. As both languages are syntactically similar, adaptations were trivial.

In TCLBG, all implementations aim to achieve the best possible performance. Therefore, the only modifications applied to those benchmarks were the ones needed to execute them in the Android environment.

All benchmarks were executed using a preset workload, individual to each benchmark. The size of each workload was determined in a way that the benchmark was executed

for at least 20s. Considering the two versions of each of the 33 benchmarks, only 5 (out of 66) presented a standard deviation greater than 18% of the mean value for the energy measurements. This indicates that, for the experimental configuration we employed, the results are stable. All JavaScript benchmarks were executed inside a wrapper originated from Apache Cordova, as a way to maximize a real simulation between a native app and a hybrid app.

B. Running the experiments

For the tests, the device used was a Nexus 5 (2013), running Android 5.1, with 16gb of flash memory, 2gb of RAM memory, chipset Qualcomm MSM8974 Snapdragon 800, CPU Quad-core 2.3 GHz Krait 400 and Li-Po 2300 mAh battery. Each benchmark and app was executed 8 times in each programming language. All data about energy consumption was collected using Project Volta. Execution time measurements were obtained using custom-built scripts.

All the experiments were executed observing the following step-by-step procedure:

- 1) Verify whether the device battery is at least 80% charged, preventing it from entering battery-save mode and keep the voltage in at least 4V.
- 2) Close all running applications not involved in the tests, activating airplane mode, and immediately rebooting the device. This step aims to isolate the application behavior, preventing other apps from remaining running or sensors such as Wi-Fi or GPS from interfering with the results.
- 3) Connect the device to the computer and reset all data of battery consumption, disconnecting it afterwards.
- 4) Execute the application or benchmark.
- 5) Prevent the app from running in the background at all times, not locking the screen, not allowing the screen to shut down or changing to another app. By locking the screen, the app starts to run in the background, not using the CPU as it normally would.
- 6) Plug the device to the computer to verify battery consumption data and execution time.

III. STUDY RESULTS

This section presents the results for the two research questions. Section III-A presents the results for RQ1 whereas Section III-B examines RQ2.

A. Is there a more energy-efficient app development approach?

In our experiments, the JavaScript versions of most benchmarks consumed less energy and were faster. However, this result was not universal and in some of the benchmarks the Java version was faster or consumed less energy.

Figure 1 shows the execution time (lines) and energy consumption (bars) of the Rosetta Code benchmarks. Overall, the JavaScript benchmarks exhibit lower energy consumption and execution time. The Java versions of these benchmarks consume a median 2.09x more energy than their JavaScript counterparts. Furthermore, the Java versions spend a median of 1.52x more time to finish the execution. The figure shows that




Fig. 1. Results of the benchmarks from Rosetta Code. The bars are sorted using the relative gain in energy consumption $\frac{\text{JavaScript energy consumption}}{\text{Java energy consumption}}$ for each benchmark.

in 18 out of 22 of the benchmarks from Rosetta, the JavaScript versions consumes less energy and in 16 they exhibited lower execution time. The results for the benchmark “sequence of non-squares” are not in Figure 1. It was omitted for the sake of graphic readability. The average time and consumption in Java were 538s and 378J and in JavaScript they were 16s e 10J, respectively. Although the result of this benchmark is not present in the graph, it was included in the calculations for consolidated results alongside all other benchmarks. Finally, in 3 of the 7 benchmarks where Java was faster, it also consumed more energy than JavaScript, which suggests a non-linear relation between energy and performance.

Figure 2 presents the results for the benchmarks from TCLBG. The Java versions of these benchmarks consume a median 1.82x more energy than their JavaScript counterparts. The figure shows that 7 out of the 10 benchmarks consume less energy in JavaScript. However, differently from the benchmarks from Rosetta Code, the median execution time of the Java versions of the benchmarks from TCLBG is 0.67x that of the JavaScript versions. Overall, 6 benchmarks are faster in Java than in JavaScript. The main reason for this behavior is the use of parallelism. Whereas all the JavaScript versions run sequentially, the Java versions of 8 benchmarks are capable of leveraging multicore processors to improve performance. 5 of these the Java version outperforms the corresponding JavaScript version. However, only two of them also exhibit lower energy consumption. This is consistent with previous work [12] that found that, for programs capable of benefiting from multi-core processors, performance is often not a proxy for energy consumption.

Analyzing all benchmarks from Rosetta and TCLGB, in the cases where JavaScript had the best results, Java applications consumed a median 2.28x more energy and their executions where a median 1.59x longer than their JavaScript counterparts. When Java performed better, JavaScript consumed a median 1.40x more energy and a median 1.40x more time.

In a preliminary study such as this, it is difficult to define general heuristics to determine which approach is better in a given situation. However, we observed that benchmarks that relied heavily on the CPU and performed many simple mathematical operations were the ones with the biggest differences in energy consumption, favoring JavaScript. Moreover, it is easier to leverage parallelism in Java, which contributes positively to the performance of the Java versions of the benchmarks. Nonetheless, that extra performance does not

seem to amount to reduced energy consumption in general. Finally, benchmarks that relied heavily on memory or files did not favor any language.

B. Can a hybrid approach to app development save energy?

The results discussed in Section III-A suggest that the two approaches for app development in Android have different trade-offs in terms of energy consumption. However, Android applications are predominantly written in Java - in a random sample of 109 apps among the 1,600 in F-Droid, only 5 employed JavaScript in any way. However, in Android, it is possible for Java code to invoke JavaScript and vice-versa. Thus, since Java is the predominant approach to write Android apps, it may be possible to save energy by retrofitting existing apps to perform part of their work in JavaScript. The major obstacle to this approach is that there is the overhead of cross-language invocations [4]. In this section we examine whether it is possible to overcome this overhead so as to make existing apps more energy-efficient, albeit perhaps harder to maintain.

Benchmarks are not useful to provide an answer for RQ2, since they work differently from apps [13]. Therefore, we have used two real-world, open-source apps for this part of the study. These apps appear in the last row of Table I. Tri Rose is an app that mathematically generates unique and intricate rose graphs and anDOF is an app to calculate depth of field for photography. These apps were chosen because, even though they spend much of their time on input and output operations, they perform a non-trivial amount of computation. TriRose comprises 1kLoC and anDOF 1.7kLoC.

When reengineering parts of a Java app to use JavaScript, it is important to determine the frequency with which the Java part will invoke the JavaScript part. If the former invokes the latter too frequently, much of the execution time and possibly energy consumption will be dominated by the overhead of cross-language invocations. If these invocations are too infrequent, application functionality may be compromised. In this work, we used three different approaches to manage cross-language invocations. In the *Stepwise* approach, one method in Java is mapped directly to a function in JavaScript and each time the method is supposed to be called, the function is used instead. In the *Batch* approach, one method in Java is mapped to a function in JavaScript, bundling several calls of the method. This function returns an aggregated result that the Java part processes to update the screen. This approach reduces the communication overhead by dividing processing




Fig. 2. Results of the benchmarks from The Computer Language Benchmark Game (TCLBG). The bars are sorted using the relative gain in energy consumption $\frac{\text{JavaScript energy consumption}}{\text{Java energy consumption}}$ for each benchmark.

duties between JavaScript and Java. Finally, in the *Export* approach, all the work to be performed in a sequence of method invocations in the original version is mapped to a single JavaScript function. Creating a hybrid app by modifying an existent one instead of coding a new app was a choice that allowed to verify whether it was possible to increase energy efficiency with minor modifications, meaning minimum effort for developers.

Tri Rose and anDOF have different behaviors. For the former, the delay of waiting for JavaScript to perform the entire computation (the *Export* approach) is not acceptable, while it is for anDOF. Thus, for anDOF we employed the *Stepwise* and *Export* approaches and for Tri Rose *Stepwise* and *Batch*. The specific workload for each app is determined in a way to try to make each execution run in approximately 30s, keeping a low relative standard deviation. The workload for the *Stepwise* approach for anDOF is 3×10^3 changes in a scroll that controls the DOF. For each change, a method is called to recalculate the depth of field. The workload for *Export* is 4×10^6 changes since, for this benchmark, it runs three orders of magnitude faster than the *Stepwise* approach. In Tri Rose, the workload for both *Stepwise* and *Batch* approaches consisted of drawing 1.5×10^3 lines on the screen, since the execution times were more similar.

Figure 3 presents the results. In both apps, using *Stepwise* degraded performance and boosted energy consumption. This indicates that the amount of overhead generated by the thousands of requests made in JavaScript increases the energy consumption to a point where it is not possible to reverse the situation with a possibly faster execution in JavaScript. This result suggests that unless it is possible to group parts of the work so as to minimize this overhead, building a hybrid app will not save energy. In every *Stepwise* test, the original Java version had a better performance and lower energy consumption.

Using *Batch*, the execution was changed to keep the results of the calculations of the points that were used to draw the curves in a buffer of 11×10^4 positions. We apply this modification to both the Java version and the hybrid version, thus producing two Java versions of this app, besides the hybrid version. The Java version using this *Batch* approach consumed 30% more energy than the hybrid version. One can note that despite the great number of data computed in JavaScript, this application spent most of the time drawing on the screen, using Java code. JavaScript had a lower energy




Fig. 3. Results from modified apps. The three approaches differ in the adaptation of the method from Java to JavaScript. *Export*: executes the method several of times, returning a single value, *Batch*: executes the method several times, returning clustered data *Stepwise*: executes the method once, several times, returning a single value each time. The bars are sorted using the relative gain in energy consumption $\frac{\text{JavaScript energy consumption}}{\text{Java energy consumption}}$ for each benchmark.

consumption but Java had a better performance.

We only apply the *Export* to anDOF. It would not make sense in the context of Tri Rose because the latter needs to continually update the screen. Updating the screen from JavaScript code in a Java app is non-trivial because Java and JavaScript employ different paradigms for user interface. The results using the *Export* approach represent a (potentially unrealistic) best case scenario, since the cross-language invocation overhead is almost entirely diluted. Nevertheless, it indicates that if an application needs to perform a substantial number of calculations, using JavaScript could lead to a significant improvement in energy efficiency and performance. The Java version consumed 35.69x more energy and took 32.77x longer.

Even though these modifications promoted considerable improvements in energy efficiency, they did not require large-scale modifications. JavaScript files for each app had 160 (anDof) and 100 (Tri Rose) LoC. Changes were relatively simple and represent less than 10% of the code of each app.

IV. THREATS TO VALIDITY

To minimize the risk that one defective device would undermine the research, all benchmarks from TCLBG were executed in a second device with similar specs, but from another manufacturer and Android version. Raw values of time and energy were similar and the relation between which of the languages performed better for each benchmark remained the same. The second device used was a Moto G3 (2015) using Android 5.1.1, with 8gb of HD, 1gb of RAM memory, chipset

Qualcomm MSM8916 Snapdragon 410, CPU Quad-core 1.4 GHz Cortex-A53 and Li-Ion 2470 mAh battery.

Project Volta is a tool for measuring energy whose accuracy has not yet been assessed. However, the most important data is the comparison and relation between the performance and energy consumption for each language, and not their raw values. Ergo, Volta's precision is not relevant to our case.

Benchmarks do not represent the behavior of an application using JavaScript as main programming language [13], and for that reason it is not possible to extrapolate the results for all applications, since applications are usually much more IO-intensive. Although this is true, benchmarks provide insight on scenarios where the performance gain is measured, by isolating usage pattern behavior. In our case, the focus was on apps that make intensive use of the CPU.

It is possible to write parts of the app in C/C++ using the Native Development Kit (NDK), which aims to improve performance. Nevertheless, due to the innate complexities of the NDK, its use is scarce, as pointed out in Google's own website⁷. It is not possible to create a app entirely with NDK. Therefore, our choice is based on industry practices.

V. RELATED WORK

Energy consumption is a hot issue in software development as a whole, not only in mobile development. One can find several papers regarding software energy consumption optimization [11], [12]. Pathak et al. [10] proposed the first fine-grained energy profiler to investigate where the energy is spent inside an app. Some have attempted to find which methods [3], API-calls [8] and applications [14] for Android are more energy-hungry. Some even specifically aim to find which source code lines are the most battery-draining [5], [6]. To the best of our knowledge, no previous work has analyzed the impact of different app development approaches on energy consumption.

The study by Charland et al. [2] is the closest work to this paper, since it compares the native and Web app development approaches. However, it focuses on user interface code, user experience, and performance for remote web apps. In particular, it does not present data on battery consumption or performance of native applications and local web applications. Nanz et. al. [9] have used the Rosetta Code tasks to compare eight different programming languages in terms of the runtime performance, and memory usage, among other factors. It does not focus on Android, however, nor analyze energy consumption.

VI. CONCLUSION

This research aimed to start an investigation on whether there is a more energy-efficient approach for android app development. It sheds some light on the strengths and weaknesses of each approach, based on experiments with benchmarks and hybridized apps that use both Java and Javascript.

⁷Google about NDK: "(...) has little value for many types of Android apps. It is often not worth the additional complexity it inevitably brings to the development process."

Preliminary results suggest that, at least for CPU-intensive operations, JavaScript outperforms Java. We have also conducted a preliminary analysis on the potential benefits of using a hybrid approach for app development, considering that most apps are written in Java. We found out that, if it is possible to avoid the need to perform multiple cross-language invocations, using a hybrid approach may lead to energy savings.

Acknowledgments. We would like to thank the anonymous reviewers for their helpful comments. Fernando is supported by CNPq/Brazil (304755/2014-1, 487549/2012-0 and 477139/2013-2), FACEPE/Brazil (APQ- 0839-1.03/14) and INES (CNPq 573964/2008-4, FACEPE APQ-1037-1.03/08, and FACEPE APQ-0388-1.03/14). Wellington is supported by FACEPE/Brazil (1052-1.03/13)

REFERENCES

- [1] Google android developer website. <https://developer.android.com/about/versions/android-5.0.html#Power>. Accessed: 2015-11-28.
- [2] A. Charland and B. Leroux. Mobile application development: Web vs. native. *Commun. ACM*, 54(5):49–53, May 2011.
- [3] M. Couto, T. Carção, J. Cunha, J. P. Fernandes, and J. Saraiva. Detecting anomalous energy consumption in android applications. In *Programming Languages - 18th Brazilian Symposium, SBLP 2014, Maceio, Brazil, October 2-3, 2014. Proceedings*, pages 77–91, 2014.
- [4] M. Grimmer, M. Rigger, L. Stadler, R. Schatz, and H. Mössenböck. An efficient native function interface for java. In *Proceedings of the 2013 International Conference on Principles and Practices of Programming on the Java Platform: Virtual Machines, Languages, and Tools*, pages 35–44. ACM, 2013.
- [5] S. Hao, D. Li, W. G. J. Halfond, and R. Govindan. Estimating mobile application energy consumption using program analysis. In *Proceedings of the 2013 International Conference on Software Engineering, ICSE '13*, pages 92–101, Piscataway, NJ, USA, 2013. IEEE Press.
- [6] D. Li, S. Hao, W. G. J. Halfond, and R. Govindan. Calculating source line level energy information for android applications. In *Proceedings of the 2013 International Symposium on Software Testing and Analysis, ISSTA 2013*, pages 78–89, New York, NY, USA, 2013. ACM.
- [7] L. G. Lima, G. Melfe, F. Soares-Neto, P. Lieuthier, J. ao Paulo Fernandes, and F. Castor. Haskell in green land: Analyzing the energy behavior of a purely functional language. In *Submitted to SANER'2016*, 2016.
- [8] M. Linares-Vázquez, G. Bavota, C. Bernal-Cárdenas, R. Oliveto, M. Di Penta, and D. Poshyvanyk. Mining energy-greedy api usage patterns in android apps: An empirical study. In *Proceedings of the 11th Working Conference on Mining Software Repositories, MSR 2014*, pages 2–11, New York, NY, USA, 2014. ACM.
- [9] S. Nanz and C. A. Furia. A comparative study of programming languages in rosetta code. In *Proceedings of the 37th International Conference on Software Engineering - Volume 1, ICSE '15*, pages 778–788, Piscataway, NJ, USA, 2015. IEEE Press.
- [10] A. Pathak, Y. C. Hu, and M. Zhang. Where is the energy spent inside my app?: Fine grained energy accounting on smartphones with eprof. In *Proceedings of the 7th ACM European Conference on Computer Systems, EuroSys '12*, pages 29–42, New York, NY, USA, 2012. ACM.
- [11] G. Pinto. Refactoring multicore applications towards energy efficiency. In *Proceedings of the 2013 Companion Publication for Conference on Systems, Programming, & Applications: Software for Humanity, SPLASH '13*, pages 61–64, New York, NY, USA, 2013. ACM.
- [12] G. Pinto, F. Castor, and Y. D. Liu. Understanding energy behaviors of thread management constructs. *SIGPLAN Not.*, 49(10):345–360, Oct. 2014.
- [13] P. Ratanaworabhan, B. Livshits, and B. G. Zorn. Jsmeter: comparing the behavior of javascript benchmarks with real web applications. In *Proceedings of the 2010 USENIX conference on Web application development*, pages 3–3, 2010.
- [14] C. Wilke, C. Piechnick, S. Richly, G. Püschel, S. Götz, and U. Assmann. Comparing mobile applications' energy consumption. In *Proceedings of the 28th Annual ACM Symposium on Applied Computing, SAC '13*, pages 1177–1179, New York, NY, USA, 2013. ACM.